

# Java

## Java Concurrency

Александр Помосов

Отметьтесь на портале

Обновите репозиторий

# Agenda



Multithreading basics



Concurrency challenges



`java.util.concurrent`



Thread-safety recipes

# Agenda



Multithreading basics



Concurrency challenges




`java.util.concurrent`



Thread-safety recipes

# Threads revisited




- Creates threads (clone syscall)
- Schedules threads (context switch)
- Provides api for Thread management

Behaviour of multithreaded program is (inter alia) dependent on **OS scheduling**

# Threads start example

The order in which threads start  
is not defined  
and is dependent on  
OS scheduling

@see  `races.RandomRunExample`





# Agenda



Multithreading basics



Concurrency challenges



`java.util.concurrent`



Thread-safety recipes

# Challenge 1. Race condition

## **Race condition** (СОСТОЯНИЕ ГОНИ, ГОНКА)

program behaviour where the output is dependent on the sequence or timing of other uncontrollable events

Parallel programs are racy by nature, some races may be errors.

@see `aces`

deadly race: <https://ru.wikipedia.org/wiki/Therac-25>

# Challenge 2. Data races

## Data race

- two or more threads in a **single process** access the same memory location concurrently, and
- at least one of the accesses is for writing

@see `data_races`

**Solution** - allow only one thread to access data at a time

## Java concurrent visualization

```
cd lecture9/
```

```
java -jar javaConcurrentAnimated.jar
```

## **Mutex (mutual exclusion)/lock**

mechanism that allow only one thread to enter 'critical section' (block of code that must be executed only by one thread at a time). That thread acquires lock

## **Reentrant lock**

lock, that can be acquired by single thread multiple times

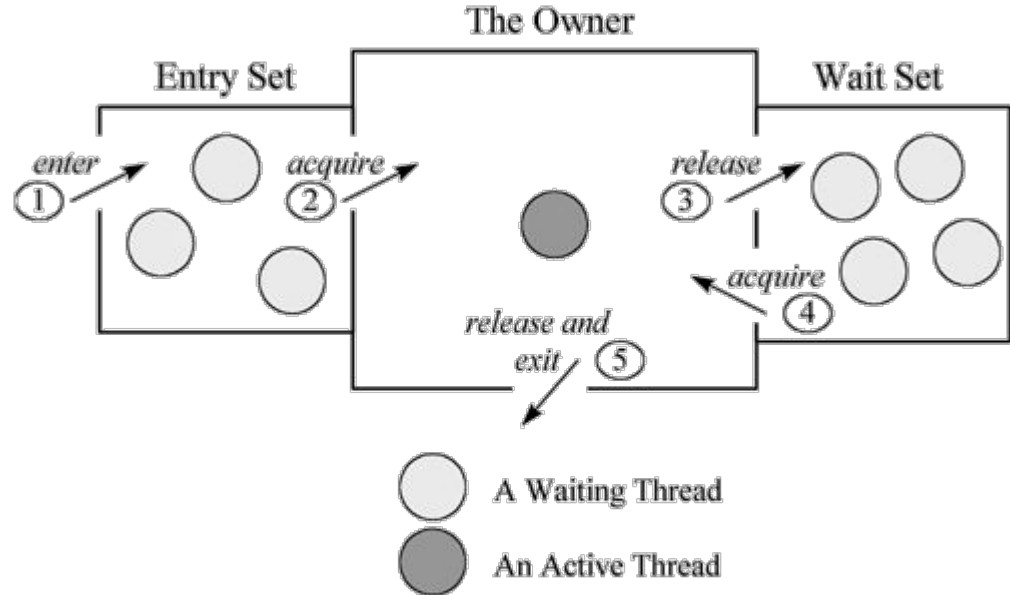
@see `javaConcurrentAnimated.jar` (`ReentrantLock`)

# Monitor

## Monitor

(mutex + entry set)

Only one thread at a time may own a monitor. Any other threads attempting to lock that monitor are blocked until they can obtain a lock on that monitor.



# Java Object internal monitors

**In java every Object has internal monitor.**

That is, every Object can act as a lock.

@see javaConcurrentAnimated.jar (synchronized)

@see synchronized\_example

Let's fix data\_races example - make increment a critical section

@see data\_races (**2 balls**)


# Challenge 3. Deadlock

## Deadlock

a state in which each member of a group of actions, is waiting for some other member to release a lock

(Do not forget **starvation** and **livelock**)

@see deadlock





# jstack - detect deadlock

**jstack** prints Java stack traces of Java threads for a given Java process or core file or a remote debug server.

<http://docs.oracle.com/javase/7/docs/technotes/tools/share/jstack.html>

**jstack** <pid>

# Java Object internal monitors

What if I want to control monitor - to control internal monitor from program?

Internal Object monitor have '**wait set**' and methods for controlling waiting threads:


**Object.wait()**, **Object.wait(timeout)** - current thread releases the monitor and enters wait set

**Object.notify()**, **Object.notifyAll()** - removes random (all) threads from wait set into blocking set

@see `java.lang.Object`

# Blocking queue example

Blocking queue blocks reading thread, when there is nothing to read. And blocks writing thread, when it is full



@see `javaConcurrentAnimated.jar`  
(`BlockingQueue`)  
@see `blocking_queue`

**Threads** (unlike processes) can communicate via **shared memory** (shared variables/shared mutable state)

If we have reads and writes to shared mutable variables, that affect each other, we call it **concurrent access**

Oracle guide to concurrency in Java


<https://docs.oracle.com/javase/tutorial/essential/concurrency/index.html>

# Thread-safety

A class is **thread-safe** if it behaves correctly when accessed from multiple threads, regardless of the scheduling or interleaving of the execution of those threads by the runtime environment, and with no additional synchronization or other coordination on the part of the calling code

(from JCiP)

<https://www.amazon.com/Java-Concurrency-Practice-Brian-Goetz/dp/0321349601>



# No state - no problem

It is **data** that must be protected from concurrent access. Even though 'synchronized' about code.

**Every access** to protected data must be synchronized (every read and write), else program is not properly synchronized.

Stateless objects are always thread-safe

So let's avoid state! (**no, impossible in practice**)

# Shared mutable state

Concurrency appear when shared mutable state is accessed from several threads

**Let's avoid concurrency!**  
(really good idea!)

Immutable and unshared objects are always **thread-safe**

@see `shared_mutable_state`

# Use immutable state (final)

‘**final**’ guarantees that after construction reference will be always read properly

- **final** only guarantees immutability for single reference to make object fully-immutable you must mark every reference final, not only root object
- if you change final fields via **reflection** - you lose guarantees

**Make final as much shared variables as possible**



# Unshared state (ThreadLocal)

```
ThreadLocal<Object> locals = new ThreadLocal<Object>();
```

As with final - ThreadLocal only guarantees, that the reference, that is accessed via ThreadLocal variable ('locals' in example) is thread local, no in-depth thread locality.

@see thread\_local

# What happen when we have SMS?

This is defined by **Java Memory Model (JMM)**

Java Language Specification (JLS)

Chapter 17. Threads and locks

<https://docs.oracle.com/javase/specs/jls/se7/html/jls-17.html>

(do not read! first look at <https://shipilev.net/#jmm> )

JMM is tricky to understand and is hard to use directly.

JMM specifies what can be read by particular read action in program.

More precisely it defines **guarantees** on read/write atomicity, write visibility and instruction ordering.

# JMM. Why so complex?

JVM is highly optimized. It is possible because of relatively **weak guarantees** of JMM.

JMM was created as a **trade-off between performance, complexity of JVM and abilities of hardware.**

JMM considered to be one of the most successful memory models.

Recently Introduced C++ Memory Model is highly based on JMM.

# Challenge 4. Atomicity

Some operations that are expected to be atomic - are not:

- `i++;`
- double/long reads and writes on 32 bit systems
- check then act actions:

```
if (!map.containsKey(key)) {  
    map.put(key, value);  
}
```

@see `data_races`


# Challenge 5. Visibility

Modern processors have multi-level caches. Thus threads running on different processors may not see changes made by other threads. It actually depends on cache coherence protocol.

[https://en.wikipedia.org/wiki/Cache\\_coherence](https://en.wikipedia.org/wiki/Cache_coherence)

Most modern processors provide coherent caches, so visibility problems are rare

**But WORA!** @see visibility



# Challenge 6. Ordering

In sake of performance javac, jit and JVM may change your code whenever it is accepted by Java Memory Model, that is reorder instructions.

After all, processor reorders instructions by himself.

**JMM restrict some reorderings.**

# Challenge 7. Performance

Reasoning about performance of concurrent programs is tricky

@see <https://shipilev.net/>



**‘volatile’** means:

- atomic **reads and writes to reference** (not all operations on object)
- reads to volatile variables always return right value
- reads and writes of volatile variables can not be reordered
- ‘happens-before’ relation

(If you make all references in your program ‘volatile’, there will be no data races)

@see `volatile_example`

# 7 challenges of concurrency

*“Я прочитал вашу лекцию  
все систематизировал и нарисовал”*

© Иероним Босх





Et sic aliter gloriā et in fine pudentia  
vultu lapidet et trahunt ad novellā pudentia



et bicondā tunc nūā ab as: et gloriādo novellā tūp



# Agenda



Multithreading basics



Concurrency challenges



`java.util.concurrent`



Thread-safety recipes

It is hard to reason low-level JMM categories, but there are a number of high-level constructions in JDK

Atomics provide non-blocking operations on common objects.  
Also provides methods for atomic '**check then act**' operations  
(compareAndSet, incrementAndGet)

@see data\_races/Stopper.java  
@see javaConcurrentAnimated.jar  
(AtomicInteger)


- ConcurrentHashMap
- BlockingQueue

@see javaConcurrentAnimated.jar  
(ConcurrentHashMap, BlockingQueue)

# ConcurrentHashMap

A hash table supporting  
**full concurrency**  
of retrievals  
and **high expected**  
**concurrency** for updates.

Iteration over **copy**  
of collection at some point




<https://habrahabr.ru/post/132884/>

<https://docs.oracle.com/javase/8/docs/api/java/util/concurrent/ConcurrentHashMap.html>



# ForkJoinPool



@see `javaConcurrentAnimated.jar`

Future is implementation of '**promises**'.

A Future represents the result of an asynchronous computation

We can block on **get()** until the result is ready.

@see `javaConcurrentAnimated.jar` (Future)

It is hard to reason low-level JMM categories, but there are a number of high-level constructions in JDK

@see `javaConcurrentAnimated.jar`  
(`CyclicBarrier`, `Phaser`)

Good manual with visualization

<https://habrahabr.ru/post/277669/>

# Agenda



Multithreading basics



Concurrency challenges



`java.util.concurrent`



Thread-safety recipes

# Recipe 1. Unshared immutable state TEXHOATOM

---

- Don't share the state variable across threads
- Make the state variable immutable
- Use synchronization whenever accessing the state variable

# Recipe 2. Safe initialization

Do not share **'this'**  
from constructor

Bad example ----->

```
public class SomeClass{  
    private Object v1;  
    private final Object v2;  
    SomeClass(Object v1, Object v2){  
        //'this' leaks  
        //initailization is not completed  
        StaticRegistry.register(this);  
        this.v1 = v1;  
        this.v2 = v2;  
    }  
}
```

@see initialization

# Recipe 3. Safe publication

To publish an object safely, **both the reference to the object and the object's state** must be made visible to other threads at the same time. A properly constructed object can be safely published by:

- Initializing an object reference from a **static initializer**
- Storing a reference to it into a **volatile field** or **AtomicReference**
- Storing a reference to it into a **final field** of a properly constructed object
- Storing a reference to it into a **field that is properly guarded by a lock**

# Recipe 4. Use JDK constructions

That to use in real life situations (by priority)

1. concurrent collections/synchronizers/ForkJoinPool
2. synchronized/volatile for reference reads/writes
3. atomics
4. wait/notify
5. volatile 'happens-before' magic

(If you are doing 4 or 5 for simple task, maybe you are doing something wrong)



# References

java concurrency in practice (signature book for Java Developer)

<https://www.amazon.com/Java-Concurrency-Practice-Brian-Goetz/dp/0321349601>

Shipilev blog (JMM, concurrency, performance, benchmarks for people, JDK contributor)

<https://shipilev.net/>

Doug Lea-s home page (java.util.concurrent father and famous spec in concurrency and allocators)

<http://g.oswego.edu/>

Java Memory Model Pragmatics (best explanation of JMM - available in russian)

<https://shipilev.net/#jmm>

JMM Under the hood (deep explanation of JMM)

<http://gvsmirnov.ru/blog/tech/2014/02/10/jmm-under-the-hood.html>

What Every Dev Must Know About Multithreaded Apps (Common knowledge)

<https://lyle.smu.edu/~coyle/cse8313/handouts.fall06/s04.msdn.multithreading.pdf>

Most active russian community on java, concurrency and related topics

<http://razbor-poletov.com/> (podcast)

<https://gitter.im/razbor-poletov/razbor-poletov.github.com> (chat)

# Reasoning about concurrent program

Bugs in concurrent programs are hard to reproduce  
Hopefully we have toolchain for analysis of multithreaded programs

jcstress

<http://openjdk.java.net/projects/code-tools/jcstress/>

(requires JDK9)

**Вот теперь понятно, спасибо!**



# Спасибо за внимание!

Александр Помосов

[alpieex@gmail.com](mailto:alpieex@gmail.com)